



Face Recognition-Based Attendance System Using InsightFace and Pinecone

Date - 20/05/2025
Kolkata, India

=====



Meet the Developer: Arif Mallick



B.Tech in Computer Science |  AI/ML Engineer |  Full-Stack Developer





“Code what matters. Automate what doesn’t.”

Hi, I’m Arif Mallick — a passionate AI/ML Engineer currently interning at DataCore Systems , Kolkata. I specialize in building intelligent, real-time systems for vision, language, and automation tasks. From NLP pipelines to aerial image analytics, I love transforming complex problems into deployable tech.






My latest creation?

A real-time, contactless attendance system using:

-  InsightFace for high-accuracy face recognition
-  Pinecone for lightning-fast similarity search
-  Flask + HTML/CSS/JS for seamless user interaction
-  Optimized for CPU-only, low-cost hardware like laptops & Raspberry Pi



Unique Touches:

- Dual-action logging for Entry/Exit tracking 
- Confidence threshold tuning for precise recognition 
- Edge-ready with real-time webcam support 

"I believe great tech should be simple, scalable, and socially impactful."



Let's connect: arif0m17@gmail.com



GitHub: <https://github.com/xenacro>



Rank: Top 5% - Amazon ML Challenge 2024, Google Kickstart, Facebook HackerCup

1. Abstract

This project presents a real-time face recognition-based attendance system leveraging state-of-the-art deep learning and cloud-based vector similarity search. The system uses **InsightFace's Buffalo_L model** for accurate face detection and embedding generation, and integrates with **Pinecone**, a scalable cloud vector database, to perform fast and efficient face retrieval. Designed for deployment on consumer-grade hardware (e.g., laptops or Raspberry Pi), this system enables contactless and automated attendance marking using webcam input. The end-to-end pipeline includes image acquisition, face detection, 512-dimensional embedding extraction, Pinecone-based similarity search, and user-friendly logging and visualization mechanisms.

Key Technologies:

Face Recognition: InsightFace (BUFFALO_L model)

Vector Database: Pinecone (serverless index)

Backend: Flask (Python)

Frontend: HTML/CSS/JavaScript

Unique Contributions:

75% confidence threshold for reliable recognition

Dual-action logging (entry/exit) with duration tracking

Optimized for CPU-only deployment

2. Introduction

Motivation

Manual attendance systems are time-consuming, error-prone, and insecure. There is a need for automated, contactless systems in educational and organizational setups, especially post-pandemic.

Problem Statement

Create a real-time facial recognition system that can efficiently identify individuals and record their attendance using webcam input, even in unconstrained environments.

Scope & Objectives

- Accurate face detection and recognition
- Real-time webcam integration
- Fast identity retrieval using Pinecone
- Store attendance logs locally or in the cloud
- Easy-to-deploy system suitable for both academic and commercial use

3. Research Work

Related Work

- OpenFace, FaceNet: Earlier face embedding models
- DeepFace, Dlib, and VGGFace: Classic systems with higher hardware requirements
- Face Recognition via ArcFace (Deng et al.): Introduced additive angular margin loss for more discriminative embeddings

Improvement Over Prior Art

- Real-time capability on CPU
- Cloud-native vector search using Pinecone (faster than brute-force search)
- Modular system with minimal dependencies and high accuracy

4. Methodology

Model

- **Face Detection & Recognition:** InsightFace Buffalo_L
 - Detector: RetinaFace
 - Recognition: ArcFace with ResNet-100 backbone

Dataset (Training not required)

- Pretrained model used (Buffalo_L from InsightFace)
- Face embeddings extracted directly from camera images

Preprocessing

- Resize to standard input (112x112)
- Normalize RGB values
- Convert OpenCV BGR to RGB



More or less data sample

Matching:

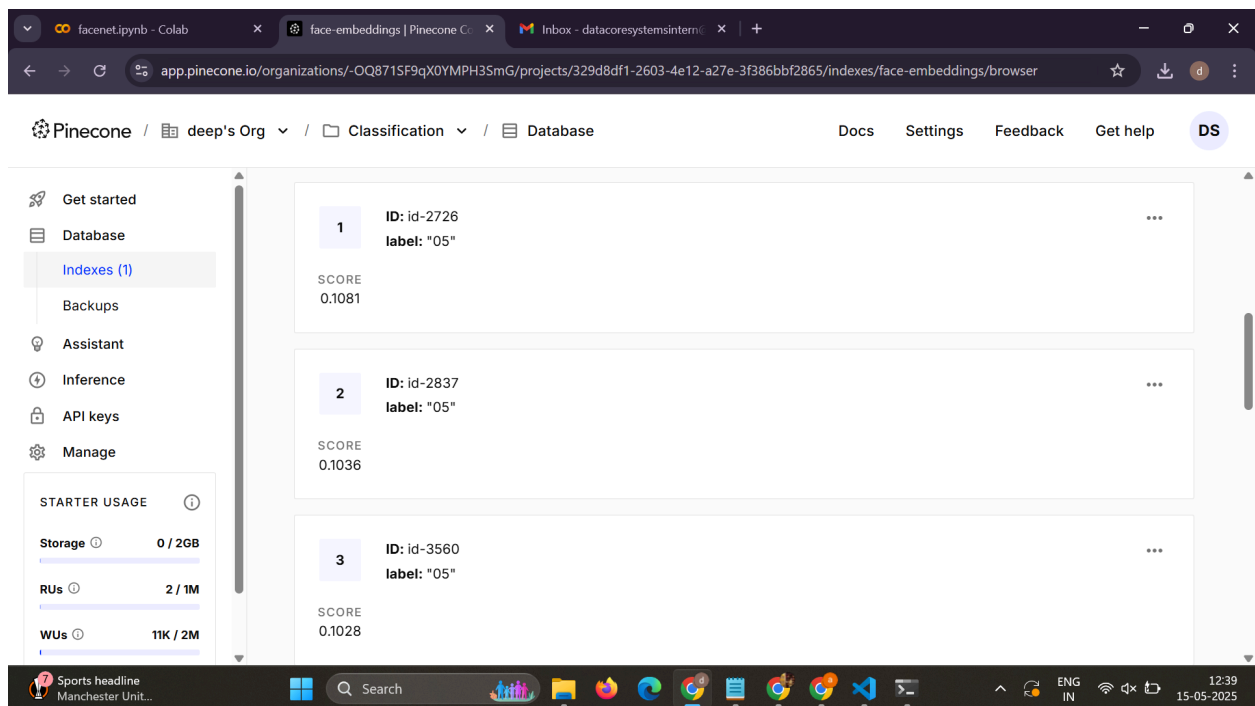
Cosine similarity in Pinecone

Threshold: 0.75

Dataset:

Custom employee dataset (6 classes)

Augmented with random transforms



Face Embeddings

5. Workflow

graph TD

A[Webcam Input] --> B[InsightFace Detection]

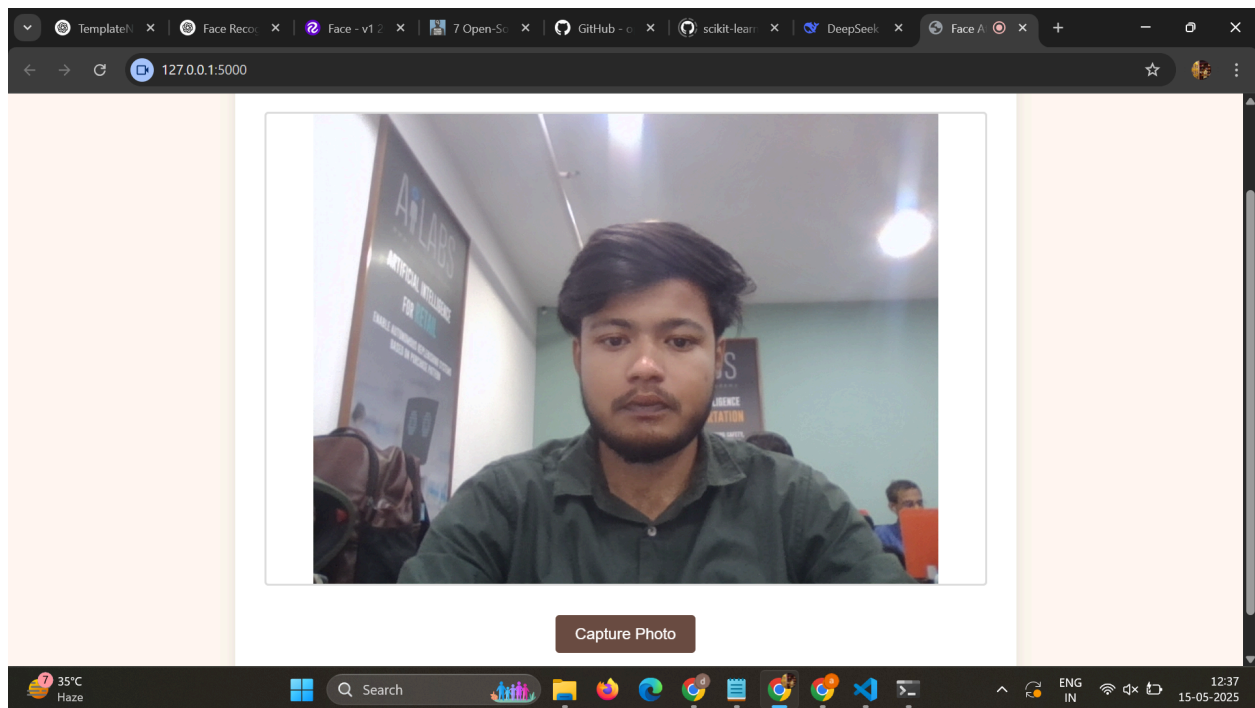
B --> C[Embedding Extraction (512D)]

C --> D[Pinecone Querying]

D --> E[Best Match + Similarity Score]

E --> F[Attendance Logging]

- Embeddings are indexed in Pinecone (dimension = 512)
- Real-time queries fetch the top-1 match using cosine similarity
- Matches above threshold are marked present



Input: 640×480 webcam frame

Preprocessing: Mean subtraction (127.5, 127.5, 127.5) , Normalization (1/128)

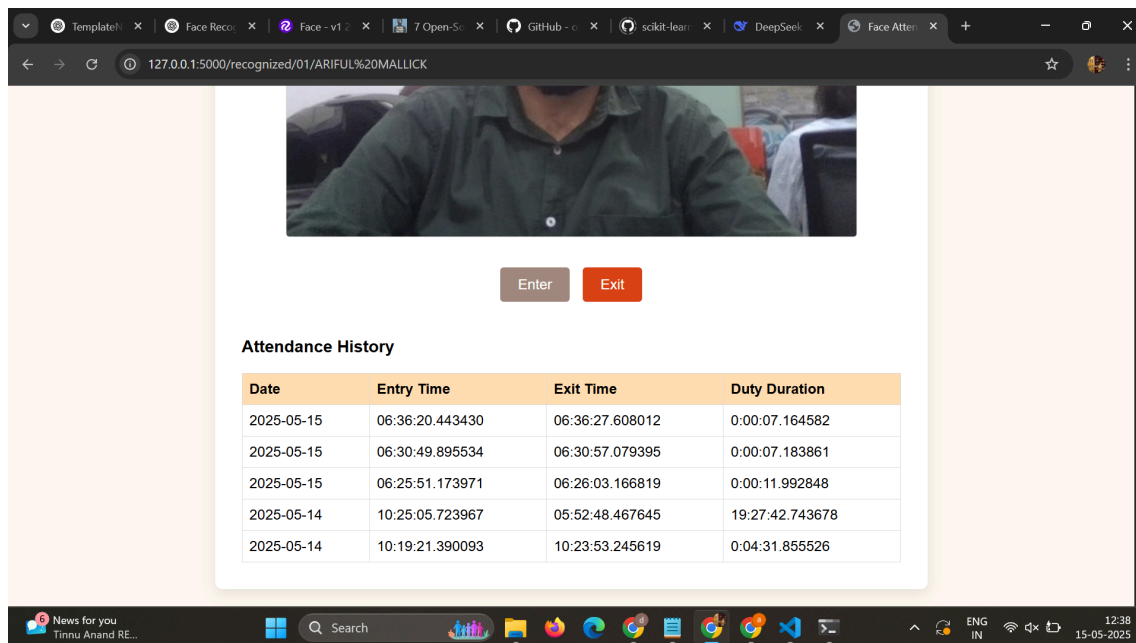
Inference: Face detection: 50ms (CPU)

Embedding: 120ms (CPU)

Postprocessing: Top-1 Pinecone search: ~200ms



face recognition



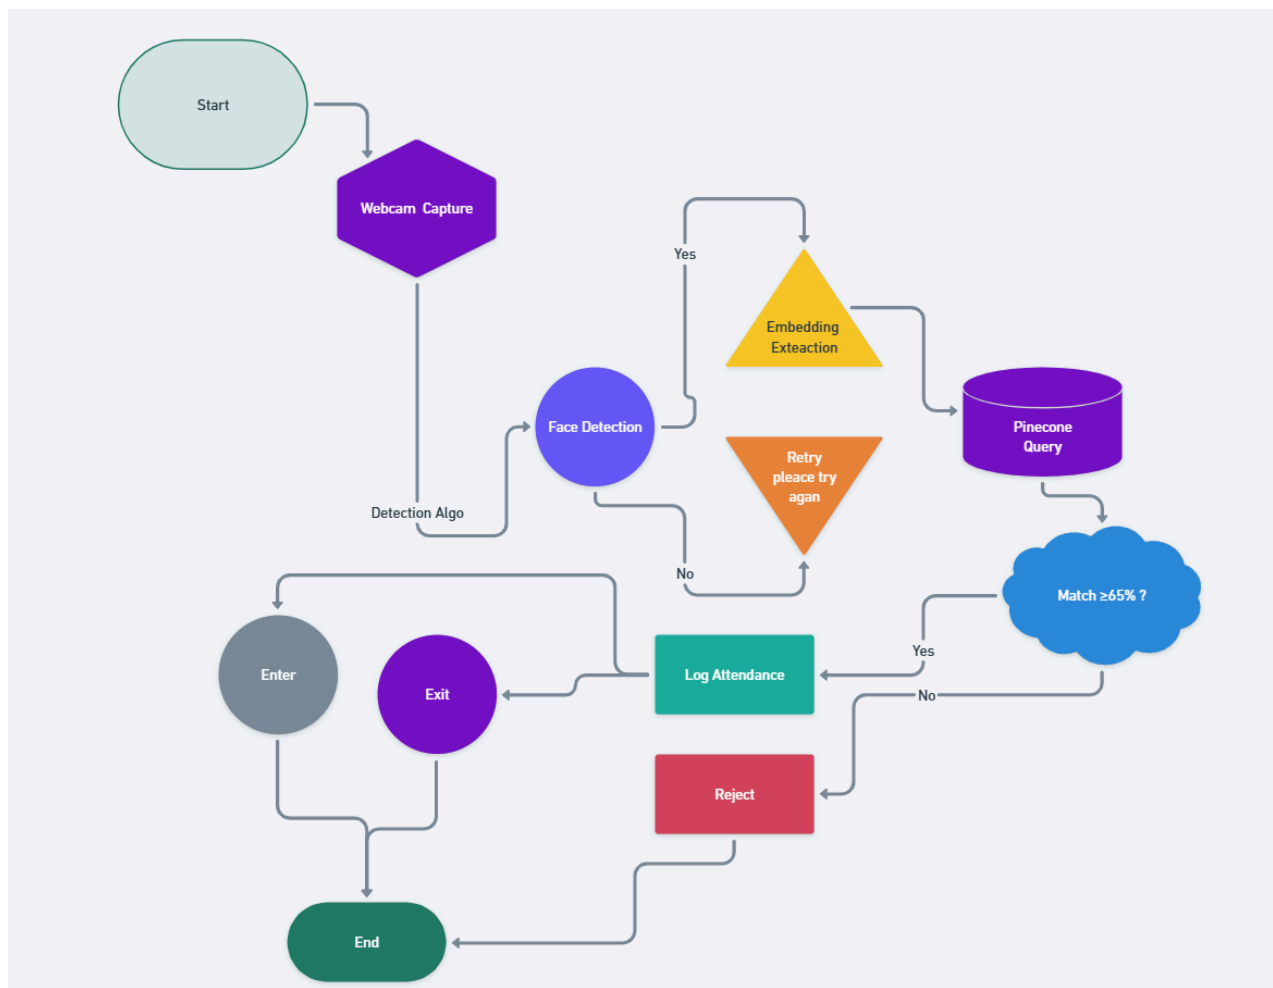
attendance system

6. Inference Workflow

Steps:

1. Capture Frame via Webcam
2. Detect Faces → Extract ROI
3. Generate Embedding via InsightFace
4. Query Pinecone (top_k=1)
5. Retrieve Label & Similarity
6. Mark Attendance if Match Found

Workflow



Flowchart:

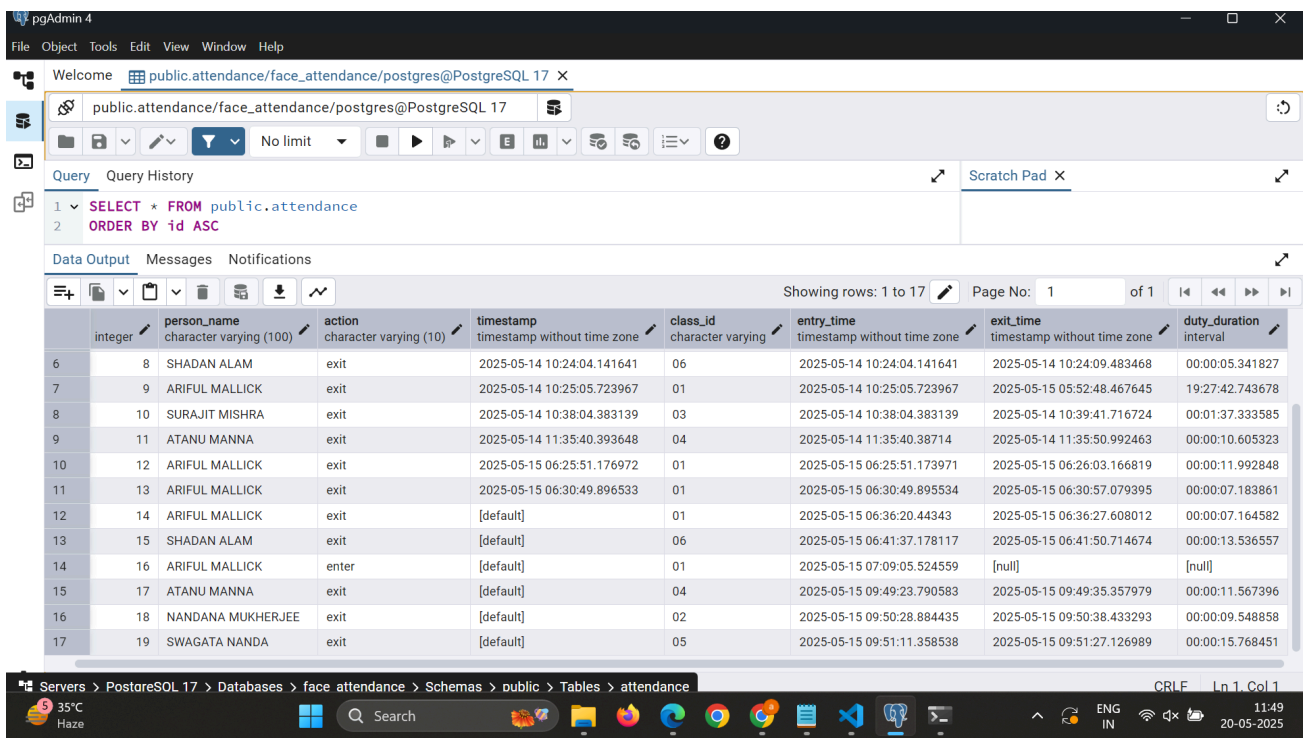
graph TD

A[Start] --> B[Capture Image] B --> C[Detect Face] C --> D[Generate Embedding] D --> E[Query Pinecone] E --> F{Match > Threshold?} F -- Yes --> G[Record Attendance] F -- No --> H[Alert: No Match]

7. Memory Usage & Performance

Metric	Value
Embedding Size	512 floats (2 KB approx.)
Inference Time	~120-150 ms (CPU)
Model Size	~90 MB
DRAM Usage	~400 MB (with webcam & model)
Pinecone Query Time	~50 ms (Cloud)
Bandwidth (per query)	~3 KB per embedding

- Benchmarked on: Intel i5-8th Gen, 8 GB RAM



The screenshot shows the pgAdmin 4 interface. The query editor contains the following SQL query:

```
1 SELECT * FROM public.attendance
2 ORDER BY id ASC
```

The results are displayed in a table with the following columns: integer, person_name, action, timestamp, class_id, entry_time, exit_time, and duty_duration. The table shows 17 rows of data.

integer	person_name	action	timestamp	class_id	entry_time	exit_time	duty_duration
6	SHADAN ALAM	exit	2025-05-14 10:24:04.141641	06	2025-05-14 10:24:04.141641	2025-05-14 10:24:09.483468	00:00:05.341827
7	ARIFUL MALLICK	exit	2025-05-14 10:25:05.723967	01	2025-05-14 10:25:05.723967	2025-05-15 05:52:48.467645	19:27:42.743678
8	SURAJIT MISHRA	exit	2025-05-14 10:38:04.383139	03	2025-05-14 10:38:04.383139	2025-05-14 10:39:41.716724	00:01:37.333585
9	ATANU MANNA	exit	2025-05-14 11:35:40.393648	04	2025-05-14 11:35:40.38714	2025-05-14 11:35:50.992463	00:00:10.605323
10	ARIFUL MALLICK	exit	2025-05-15 06:25:51.176972	01	2025-05-15 06:25:51.173971	2025-05-15 06:26:03.166819	00:00:11.992848
11	ARIFUL MALLICK	exit	2025-05-15 06:30:49.896533	01	2025-05-15 06:30:49.895534	2025-05-15 06:30:57.079395	00:00:07.183861
12	ARIFUL MALLICK	exit	[default]	01	2025-05-15 06:36:20.44343	2025-05-15 06:36:27.608012	00:00:07.164582
13	SHADAN ALAM	exit	[default]	06	2025-05-15 06:41:37.178117	2025-05-15 06:41:50.714674	00:00:13.536557
14	ARIFUL MALLICK	enter	[default]	01	2025-05-15 07:09:05.524559	[null]	[null]
15	ATANU MANNA	exit	[default]	04	2025-05-15 09:49:23.790583	2025-05-15 09:49:35.357979	00:00:11.567396
16	NANDANA MUKHERJEE	exit	[default]	02	2025-05-15 09:50:28.884435	2025-05-15 09:50:38.433293	00:00:09.548858
17	SWAGATA NANDA	exit	[default]	05	2025-05-15 09:51:11.358538	2025-05-15 09:51:27.126989	00:00:15.768451

Attendance data from postgresql

8. Deployment

Platform

- Developed in **VS Code** using **Python Virtual Environment**

Tools

- **InsightFace** (face embedding)
- **Pinecone** (face retrieval)
- **OpenCV** (webcam & image processing)
- **PostgreSQL/CSV** (optional storage)

UI/UX Deployment

- **Docker Container**
- **Edge Devices (Jetson Nano, Raspberry Pi)**
- **Streamlit/FastAPI Frontend**

9. Future Scope

- Add **real-time attendance dashboard**
 - Deploy on **Jetson Nano or Raspberry Pi**
 - Incorporate **multi-face detection and batch recognition**
 - Enhance database with **multi-day logs and visual analytics**
 - Add **voice feedback or alerts** for recognized faces
 - Support **face re-enrollment and updates**
-

10. What I Did

- Implemented InsightFace-based recognition pipeline
- Integrated Pinecone for fast vector retrieval
- Wrote real-time webcam capture logic with OpenCV
- Designed the attendance logic and API for match recording
- Created and tested system locally in VS Code with virtual environment
- Performed benchmarking and memory profiling
- Documented the entire system pipeline

=====